

# Yulia Volkova

---

Senior Data Scientist

📍 London, UK | [LinkedIn](#) | +447450514095 | [yuulia.volkoval@gmail.com](mailto:yuulia.volkoval@gmail.com)

*Senior Data Scientist with 8 years of experience delivering scalable ML solutions and NLP systems in production. Expertise in LLM evaluation & post-training, statistical modelling, and MLOps/infra. Passionate about AI alignment and technical safety, with a focus on mechanistic interpretability and AI control.*

---

## AI Alignment & Technical Safety

**Visiting Member at [London Initiative for Safe AI \(LISA\)](#)**

**Sept 2025 - present**

**SPAR Fellowship**

**Oct 2025 - present**

- CoT monitorability project under Qiyao Wei (MATS)

**[Algoverse AI Safety Research Fellowship](#)**

**Jul–Aug 2025**

- Two concurrent research projects:
  1. Mechanistic interpretability under Nicky P
  2. Cooperative AI under Max Kleiman-Weiner

**[ARBOx2 Oxford AI Safety Bootcamp](#)**

**Jul 2025**

- Participated and completed the ARENA curriculum.
- 

## Work Experience – Behavox

*Behavox is a security software company specialising in communication surveillance. Leading AI compliance solution in the marketplace.*

**Data Science Production Team Lead**

**2024–2025**

- Led development and deployment of LLM-driven Java applications across 3 products and 14 languages, using a pipeline of fine-tuned models (SBERT -> DeBERTa -> Qwen-14B)
- Managed execution of Java applications on the Apache Beam Dataflow runner; monitored GPU utilization and used BigQuery datasets to debug pipeline behavior and resolve incidents
- Led cross-team roadmapping for ML features under production constraints
- Built reusable tooling—CI/CD templates and MLflow tracking—enabling rapid MVPs for client deliveries

**Insider Crime Product Lead**

**2020–2024**

- Led R&D and production of high-recall (81% at 0.5% FPR) misconduct detection systems using RoBERTa + SBERT models architecture pipeline
- Specialised in working with LLMs, fine-tuning + pre-training them to meet Insider Crime product requirements
- Validated and proved that shorter text sequences in RoBERTa training maintain performance, doubling training speed and halving company training costs
- Built and operated a chatbot for internal workflows hosted on AWS and containerized with Docker, orchestrated via Nomad using the Slack API—serving 200+ daily users
- Developed a hybrid recommendation system leveraging RoBERTa and SBERT for text embeddings, K-Means for clustering, and integrated user attributes into a CatBoost model for personalised predictions

## Data Scientist

2017–2020

*Key responsibilities:* development of multilingual applications and reports to flag statistical deviations and compliance breaches in client data, designing and building infrastructure to support application creation, QA and maintenance

- Landed proof of concepts that resulted in 3+ years client contracts
  - Improved compliance model recalls from 70% to 90%
  - Reduced client bug requests from 1/day to 1/month by setting up an efficient QA process
  - Self-studied Java to remove the Analytics dependency on Backend teams
  - Developed best practices for code quality as one of the first Data Science hires and created foundational documentation during company scaling
- 

## Education

### MSc Economics, London School of Economics

2017–2018

— Distinction (dissertation); Advanced Macroeconomics · Econometrics · Quantitative Methods

### BSc Philosophy & Economics, London School of Economics

2013–2016

— Distinction; Mathematical Methods · Statistics · Econometrics · Philosophy of Science

---

## Skills Summary

Category	Skills
<b>Languages &amp; Tools</b>	Python, Pandas, NumPy, Scikit-learn, FastAPI, spaCy, SQL, Java, Kotlin, Groovy, Jenkins
<b>Machine Learning &amp; NLP</b>	LLM fine-tuning (LoRA/qLoRA), PyTorch, HuggingFace, prompt engineering, transformers, SBERT, PCA, t-SNE, models evaluation (precision, recall, ROC/PR-AUC, F1), Weights & Biases
<b>Infrastructure</b>	Testing automation, synthetic data generation, containerised deployment, AWS S3, GCP, Docker, Nomad, Kubernetes, TeamCity, Gitlab CI/CD pipelines

Category	Skills
<b>Statistics &amp; Econometrics</b>	Hypothesis testing, time-series analysis, A/B testing, IV estimation, panel regression, fixed/random effects, vector auto-regressions
<b>Leadership</b>	Mentorship, technical interviewing, product roadmapping, cross-team communication, internal presentations